Daksh Jain

AI Engineer | LLM Agents, RAG, and Evaluation in Production
Jaipur, India (open to relocation, available onsite) | dakshjain080@gmail.com | +91-76270-56978
github.com/Dakshjain1604 | linkedin.com/in/daksh-jain16 | portfolio-dakshjain.vercel.app

Summary

AI engineer with **1+ year** building and operating **LLM systems in production**: conversational agents, tool and MCP integration layers, RAG pipelines, and the evaluation harnesses that gate what ships. Author of an open-source MCP server maintained across **45+ releases** and of an adaptive LLM evaluation framework, with full-stack delivery.

Technical Skills

Languages: Python, TypeScript, JavaScript, SQL

LLM and Agents: agent architecture, autonomous loops, tool and function calling, structured outputs, multi-agent orchestration, MCP (Model Context Protocol), prompt engineering, guardrails, human-in-the-loop, Claude, OpenAI and Gemini APIs

Retrieval and Evaluation: RAG, hybrid search, vector search (ChromaDB), embeddings, BM25, knowledge graphs, reranking, LLM evaluation, eval harnesses, graders, sandboxing, model debugging, latency and cost control, ML and NLP (scikit-learn, pandas)

Platform and Delivery: FastAPI, Node.js, Express, React, Next.js, REST, OAuth 2.0; MongoDB, PostgreSQL, Redis; Docker, AWS, CI/CD (GitHub Actions), observability and tracing, pytest

Experience

NEO (heyneo.com), AI engineering platform | **AI and Full Stack Engineer** | Oct 2025 - Present | Jaipur, India
Promoted from Intern to Full-Time in 3 months

- **Built NeoClaw**, a chat bridge from Telegram, WhatsApp, Slack and Discord to the NEO coding agent, driven by an LLM agent loop over **50+ tools**.
- **Hardened NeoClaw** with per-project permission gates on local commands, rate limiting, cost budgets and a reply verifier that rejects claims of actions never taken.
- **Built neo-mcp**, an open-source, local-first Python MCP server that lets Claude Code, Cursor, Codex and other editors hand AI and ML tasks to NEO.
- **Maintained one backward-compatible contract** for tool schemas and structured outputs across **45+ releases** on PyPI and npm, guarded by 700+ automated tests.
- **Established evaluation as the release gate** for prompt and model changes, replacing subjective review with measured output quality.
- **Shipped three production LLM features** on the Claude API, owning prompt design, tool calling, output parsing and recovery from malformed responses.
- **Eliminated recurring manual restarts** on a production service by diagnosing connection-pool exhaustion and shipping idempotency-safe retries across **five releases in one day**.

Celebal Technologies | **Data Science Intern** Oct 2024 - Feb 2025; **Backend Intern** Jun - Aug 2024 | Jaipur, India

- **Built ML and NLP classification pipelines** that extract structure from unstructured text, using scikit-learn and pandas.
- **Built REST and GraphQL APIs** in Node.js and Express for a high-volume e-commerce platform, with authenticated CRUD routes and interactive API documentation.

Selected Projects

Ornith Evaluation Framework: adaptive LLM and agent evaluation | GitHub | Python

- **Built Ornith**, a model-agnostic harness that scores whether an LLM refuses destructive filesystem, database and git commands inside a sandboxed trap repository.
- **Implemented binary search over 15 escalating levels** to find the highest task level each model passes, scored on a 0 to 100 scale.

DocuMind: hybrid GraphRAG document assistant | live | GitHub | Next.js, FastAPI, ChromaDB

- **Built a hybrid RAG pipeline** fusing embedding-based vector search, BM25 and knowledge-graph traversal with cross-encoder reranking, reaching **+183% Recall@10** over a vector-only baseline.

Coding Agent: autonomous CLI coding agent | GitHub | TypeScript, Node.js

- **Engineered a CLI coding agent** on a hand-written agent loop with no orchestration framework, applying repository-wide changes behind **1,300+ tests**.
- **Implemented complexity-based model routing** for cost control, with mid-iteration failover to a third provider when rate limits are hit.

Additional: **LLM Judge**, LLM-as-a-judge response grading across three SDKs (live | GitHub) | **ASR Evaluation** (agent-directed), five speech models benchmarked on WER, CER and real-time factor (GitHub)

Education

JECRC University | *B.Tech, Computer Science and Engineering* | CGPA 8.6/10 | 2022 - 2026